

Overview of Transaction Processing

Outline of Lecture

- Introduction to Transaction Processing
 - Review
 - Problems and Solution
- Concepts
- Properties
- Example – Oracle DB

Databases - multi user applications

- Databases - multi user applications
- Concurrent access to the same data may induce different types of problems
- Errors may occur in the DBMS or its environment

Transactions, Recovery and Concurrency Control

- Examples of problems: hard disk failure, application/DBMS crash, division by 0, ...
- **Recovery**: activity of ensuring that, whichever of the problems occurred, the database is returned to a consistent state without any data loss afterwards
- **Concurrency control**: coordination of transactions that execute simultaneously on the same data so that they do not cause inconsistencies in the data because of mutual

Types of Failures

- **Transaction failure** results from an error in the logic that drives the transaction's operations and/or in the application logic
- **System failure** occurs if the operating system or the database system crashes
- **Media failure** occurs if the secondary storage is damaged or inaccessible

Transactions, Recovery and Concurrency Control

- Transaction: set of DML operations induced by a single user or application, that should be considered as one undividable unit of work
 - Transfer between two bank accounts
- Transaction - 'succeeds' or 'fails' in its entirety
- Transaction - database from one consistent state into another one

DML

■ Main operations

- INSERT
- UPDATE
- DELETE

SQL - DML

■ INSERT

```
INSERT INTO R (A1, A2, ..., An)  
VALUES (v1, v2, ..., vn)
```

■ UPDATE

```
UPDATE R SET <new value  
assignment>
```

```
WHERE <condition>
```

• DELETE

```
DELETE FROM R WHERE <condition>
```



Real work

- It's all about fast query response time and correctness
- DBMS is a multi-user systems
 - Many different requests
 - Some against same data items

Problems

- Failure occurs
- Concurrent database access
- Resilience to system failures


Concurrent DB access



Update PIB_account
Set money = money - 500
Where cID = 1235817

Update PIB_account
Set money = money + 759
Where cID = 1235817

Bulk Inserts



```
INSERT INTO Studio(name)
SELECT DISTINCT
studioName
FROM Movie
WHERE studioName NOT
IN
(SELECT Name
FROM Studio);
```

Failure

- Example: *Transfer an amount of money* between two accounts
 - Checks that the first has at least that much money
 - If so, transfers money from first account to second
- What if failure occurs between two updates?
- Database will be left in inconsistent state
- See that certain operations needs to be done *atomically*
 - Either all operations complete or none goes through

Solution

Solution to previous problems is to group database operations into *transactions*

- **A transaction is a sequence of one or more (DML) SQL operations treated as a unit**

ACID Transactions

A DBMS is expected to support "ACID transactions", which are:

- **Atomic:** Either the whole transaction is run, or nothing
- **Consistent:** Database constraints are preserved.
- **Isolated:** Different transactions may not interact with each other.
- **Durable:** Effects of a transaction are not lost in case of a system crash.

System Log

- Remember, DBMS must assure that we don't lose information due to system crashes
 - i.e., how do we recover from failure?
- Keep system log
 - Kept on disk, backed up periodically
 - Record every action

Controlling transactions

We can explicitly start transactions using the `START TRANSACTION` statement, and end them using `COMMIT` or `ROLLBACK`:

- `COMMIT` causes an SQL transaction to complete successfully.
 - Any modifications done by the transaction are now permanent in the database.
- `ROLLBACK` causes an SQL transaction to end by aborting it.
 - Any modifications to the database must be undone.
 - Rollbacks could be caused implicitly by errors.